

Neet Mehulkumar Mehta

India | neetmehta19@gmail.com | +91 9484508208 | GitHub: <https://github.com/neetmehta>
Website: neetmehta.github.io

MACHINE LEARNING ENGINEER

Perception AI Engineer with 2+ years of experience building production-grade autonomous driving perception systems, specializing in multi-sensor fusion, BEV-based 3D object detection, and embedded deployment. Proven track record of improving model accuracy and latency through geometry-aware learning, CUDA/TensorRT optimization, and end-to-end MLOps pipelines, and shipping perception stacks to vehicles under SIL/HIL and V-model processes. Recently applied image-conditioned generative vision models as an extension of perception expertise, focusing on spatial alignment, consistency, and controllable generation, without shifting away from core perception and robotics fundamentals.

EDUCATION

Worcester Polytechnic Institute (WPI)
Master of Science- Robotics Engineering, GPA- 3.85/4.00
Nirma University
Bachelor's in mechanical engineering, GPA- 7.8/10.00

Worcester, MA
Dec 2022
Ahmedabad, India
May 2020

KEY SKILLS

- **Programming Skills:** C++, CUDA, Python, MATLAB
- **Tools and Libraries:** ROS2, Pytorch, TensorRT, ONNX, CARLA simulator, AWS, PCL (Point Cloud Library), OpenCV, Docker, Git, ComfyUI, Blender 3D, Solidworks.

WORK EXPERIENCE

Aavaran
Applied Generative Vision Engineer, Self Employed
Pytorch, ComfyUI, Docker, Vastai

Bhuj, India
Mar 2025 – Present

Applied generative vision techniques as an extension of my perception and geometry background, focusing on spatial alignment, scale consistency, and deployment efficiency.

- Designed and deployed image-conditioned generative vision pipelines for interior visualization and virtual try-on, emphasizing geometry consistency, spatial alignment, and controllable generation.
- Built text- and image-conditioned scene editing workflows using Qwen 2.5.11 Image Edit for realistic in-context placement of home textiles with preserved scale, occlusion, and lighting.
- Created an AI brand character using FLUX.2-dev, training a LoRA for identity consistency across poses and viewpoints, and applied ControlNet for pose and composition control in advertising imagery.
- Developed a virtual try-on system for traditional apparel, transferring garments onto real subjects and AI characters while preserving fabric texture, drape, and pose alignment.
- Extended image pipelines toward short-form generative video using WAN 2.2, focusing on temporal coherence.
- Optimized inference using quantized GGUF / FP8 models on RTX 4090 (Vast.ai) and consumer GPUs via custom ComfyUI workflows.

TORC Robotics
ML Engineer II
C++, CUDA, ROS2, Python, Pytorch, AWS, TensorRT

Austin, TX
Feb 2023 – Feb 2025

- Served as SCRUM Master for the perception team, driving sprint planning, backlog grooming, cross-team coordination, and delivery tracking, while ensuring alignment between research, systems, and deployment milestones for a complex multi-sensor autonomy stack.

Learned Sensor Fusion

- Defined clear upstream and subsystem interface requirements, specifying latency, synchronization, and data quality constraints for image streams (ISP), LiDAR point cloud aggregation, and depth map inputs, translating system-level constraints into well-scoped component requirements.
- Conducted a POC evaluation of learned multi-sensor object detectors (including BEVFusion and related architectures), ultimately shortlisting BEVFusion as the baseline for further architectural innovation and production feasibility.
- Implemented depth-aware BEV unpooling, generating dense 3D camera points and features from monocular depth outputs, significantly improving camera-to-BEV feature richness and fusion quality.
- Designed and implemented an attention-based learned spatio-temporal feature aligner to jointly align LiDAR BEV features and dense 3D camera features across multiple sensors and time steps, achieving a +2.5% mAP improvement over vanilla BEVFusion on proprietary datasets.
- Built a robust model conversion and deployment pipeline, converting models from PyTorch → ONNX → TensorRT, enabling production-grade inference on embedded platforms.

- Applied **post-training quantization (PTQ) to FP16**, reducing end-to-end inference latency by ~50% while preserving detection accuracy.
- Designed and implemented a **custom TensorRT BEV pooling plugin** capable of efficiently pooling ~1M camera points into a **3D BEV grid**, addressing a critical performance bottleneck.
- Developed **custom CUDA kernels** for performance-critical operations including **voxelization, Non-Maximum Suppression (NMS), and bounding box decoding**, demonstrating deep expertise in low-level GPU optimization.
- Integrated the optimized TensorRT model into a **proprietary Inference Engine**, creating a clean abstraction layer for deployment across embedded hardware targets.
- Architected a **highly configurable ROS2 perception node** with **sensor-level abstraction**, allowing camera and LiDAR vendors to be swapped without impacting internal data formats; the node supports **arbitrary numbers of cameras and LiDARs** and is extensible to **radar integration** via modular data converters.
- Delivered outputs in **team-aligned interface formats** defined through **ICP documents**, enabling smooth downstream consumption by planning and control teams.
- Executed **full system integration and deployment on a production truck**, following a rigorous **V-model development process** including **unit testing, SIL, and HIL validation**.
- Implemented **rich perception visualization tools** using **RViz and Foxglove**, enabling efficient debugging, performance analysis, and cross-team collaboration.
- Built a **fully automated, end-to-end model lifecycle and MLOps pipeline** spanning **training → versioning → conversion → validation → deployment**, enabling seamless onboarding of newly trained models or minor architectural changes; models were **versioned and tracked in a centralized registry using Comet**, automatically pulled with all required artifacts, converted **PyTorch → ONNX → TensorRT** for embedded compute, validated through **unit, component, and performance tests**, and deployed via **CI/CD**, requiring only a **model version update in the ROS2 node** to trigger conversion, validation, and metric generation for the new model.

Domain Adaptation

- *Implemented a Proof-of-Concept (POC) for Unsupervised Domain Adaptation (UDA) based on the GUDA (Geometric Unsupervised Domain Adaptation) framework to address sim-to-real domain gaps in autonomous driving perception models.*
- *Utilized large-scale synthetic datasets (thousands of images) provided by Parallel Domain, enabling geometry-aware adaptation without requiring labeled real-world data.*
- *Designed and trained a geometry-guided adaptation pipeline for semantic segmentation and monocular depth estimation, enforcing cross-domain consistency using depth, camera pose, and geometric alignment constraints instead of pixel-level supervision.*

Monocular Depth Estimation

- Developed a state-of-the-art unsupervised monocular depth estimation model for autonomous driving scenarios, leveraging view synthesis and photometric consistency losses to learn depth without ground-truth supervision.
- Designed and implemented a self-supervised training pipeline using monocular video sequences, jointly optimizing depth and ego-motion through differentiable image warping and multi-scale reconstruction losses.
- Incorporated geometric constraints including smoothness regularization, edge-aware depth priors, and occlusion masking to improve depth stability and reduce artifacts in dynamic driving environments.
- Achieved robust depth prediction performance across diverse lighting and weather conditions by applying temporal consistency losses, data augmentation strategies, and scale-invariant normalization techniques.
- Evaluated adapted models on real-world driving data, demonstrating **improved generalization for both segmentation and depth estimation** compared to non-adapted baselines, validating the feasibility of GUDA-based sim-to-real adaptation.

TORC Robotics

Perception Engineer – Co-Op

C++, Python, Pytorch, AWS, TensorRT

- Worked on Data extraction and data postprocessing for deep learning architectures. Established extendable pipeline to generate detailed metrics report for each Deep learning model.
- Developed automated hyperparameter tuning stage in AWS Sagemaker. Used Bayesian search to find optimal hyperparameters.

Blacksburg, VA

Jan 2022 – Aug 2022

PROJECTS

Synthetic Data Generation for Self Driving Cars using CARLA simulator	Aug 2022 – Present
<i>Python, Pytorch, CARLA</i>	
• Designed and developed a scalable synthetic data generation toolkit using CARLA , enabling unlimited dataset creation for autonomous driving perception tasks including depth estimation, optical flow, instance segmentation, semantic segmentation, 2D object detection, and 3D object detection .	
• Implemented a highly configurable sensor simulation framework , allowing dynamic modification of sensor placement, intrinsics/extrinsics, field-of-view, and update rates , supporting realistic multi-sensor perception experiments.	

- Built controls to **programmatically vary scene complexity**, including **traffic density, pedestrian count, and agent behavior**, enabling systematic stress-testing of perception models under diverse operating conditions.
- Enabled **configurable data capture frequency and synchronization**, supporting fine-grained control over temporal resolution for video-based and multi-frame learning pipelines.
- Generated **multi-modal, time-synchronized annotations** across tasks, facilitating research in **sensor fusion, self-supervised learning, domain adaptation, and sim-to-real transfer**.
- Designed the repository with **modular, extensible architecture**, allowing easy integration of new sensors, annotation types, and export formats for downstream training pipelines.

June 2022 – July 2022

Self-Supervised Monocular Depth Estimation (Monodepth2) from scratch

Python, Pytorch

- Implemented **Monodepth2 from scratch in PyTorch** to gain a deep understanding of **self-supervised monocular depth estimation**, including view synthesis, photometric reconstruction, and scale-invariant depth learning.
- Reproduced the full **training pipeline** using monocular video sequences, implementing **pose estimation, differentiable image warping, multi-scale supervision, and automasking** as described in the original paper.
- Developed and integrated key loss components including **photometric reprojection loss, edge-aware smoothness regularization, and minimum reprojection selection** to handle dynamic scenes and occlusions.
- Validated implementation correctness through **qualitative depth visualization and quantitative evaluation** on standard driving datasets, ensuring consistency with reported Monodepth2 behavior.

3D Object Detection in Point Cloud using Voxelnet

Sept 2021 – Dec 2021

Python, Pytorch, OpenCV

- Implement a 3D detection network (VoxelNet) on the KITTI vision (Point Cloud) benchmark dataset to unify feature extraction and bounding box prediction into a single-stage, end-to-end trainable deep network.

Multinet-2: A Multitask learning architecture for Semantic, Depth, and Normal prediction

Feb 2022 – May 2022

Python, Pytorch

- Implemented Deep CNN architecture that can predict Semantic mask, estimate Depth, and normal simultaneously.
- Increased combined inference speed to 1.75x with a slight accuracy drop.

“Attention Is All You Need” – Transformer Reimplementation (Learning Project)

Apr 2025– May 2025

Python, Pytorch

- Reimplemented the **Transformer (“Attention Is All You Need”) from scratch in PyTorch** for English–German translation to gain a deep understanding of self-attention and encoder–decoder architectures; validated results qualitatively.

BERT from Scratch

Jun 2025 – Present

Python, Pytorch, OpenCV

- Reimplementing **BERT from scratch in PyTorch** to deeply understand **Transformer-based language models**, including token embeddings, multi-head self-attention, encoder blocks, and masked language modeling.
- Developing a **financial sentiment analysis pipeline** on top of the implementation, fine-tuning the model for domain-specific text classification.
- Validating implementation through **qualitative analysis and basic quantitative checks**, focusing on architectural correctness and training dynamics.

Real-time hand gesture recognition using SSD-MobileNet and Transfer Learning

Oct 2021– Dec 2021

Python, Tensorflow, OpenCV

- Trained object detection model consisting of 5gestures by Transfer Learning to a pre-trained SSD-MobileNet model and TensorFlow object detection API on RTX 2060 MAX-Q GPU.
- Achieved 80% accuracy for a class.
- Trained lightweight model suitable for real-time hand gesture recognition.

Popular CNN architectures

Jan 2022 – Jan 2023

Python, Pytorch

- Implementing popular Deep Learning architecture for 2D/3D object Detection, Semantic and Instance Segmentation, and Depth Estimation.
- The purpose of these projects is to develop a strong foundation of the theoretical and practical aspects of Deep Learning.
- You can find all the projects on my GitHub. Some of them might still be in development.

Implementation and Visualization of Autonomous Robot Path Planning Algorithms

Feb 2021 – May 2021

Python

- Implemented discrete and sampling-based algorithms such as A*, Weighted A*, Dijkstra, Probabilistic Road Map (PRM), Rapidly exploring Random Tree (RRT), RRT*, and Informed RRT* to navigate through obstacles in a 2D environment.